

# jwkfecq15

November 13, 2025

## 1 Insurance Charges Dataset

```
[46]: import pandas as pd

# Load the CSV file into a pandas DataFrame
df = pd.read_csv('/content/insurance[1].csv')
# Display the first 5 rows of the DataFrame to inspect the data
df.head()
```

```
[46]:   age      sex      bmi  children smoker      region      charges
0    19  female  27.900          0     yes  southwest  16884.92400
1    18     male  33.770          1     no  southeast   1725.55230
2    28     male  33.000          3     no  southeast   4449.46200
3    33     male  22.705          0     no northwest  21984.47061
4    32     male  28.880          0     no northwest   3866.85520
```

```
[47]: df.duplicated().sum()
```

```
[47]: np.int64(1)
```

```
[48]: print("DataFrame Info:")
df.info()

print("\nMissing Values:")
print(df.isnull().sum())

print("\nDescriptive Statistics:")
print(df.describe())
```

```
DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age         1338 non-null   int64  
 1   sex         1338 non-null   object  
 2   bmi         1338 non-null   float64
```

```
3    children   1338 non-null   int64
4    smoker     1338 non-null   object
5    region     1338 non-null   object
6    charges    1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

Missing Values:

```
age        0
sex        0
bmi        0
children   0
smoker     0
region     0
charges    0
dtype: int64
```

Descriptive Statistics:

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
[49]: # Remove duplicate rows
df.drop_duplicates(inplace=True)

# Verify that duplicates have been removed
print(f"Number of duplicate rows after removal: {df.duplicated().sum()}")
print(f"New shape of the DataFrame: {df.shape}")
```

Number of duplicate rows after removal: 0

New shape of the DataFrame: (1337, 7)

```
[50]: # Select 'age' as the independent variable (X) and reshape it for scikit-learn
X = df['age'].values.reshape(-1, 1)
# Select 'charges' as the dependent variable (y)
y = df['charges'].values

from sklearn.model_selection import train_test_split

# Split the data into training and testing sets (80% train, 20% test)
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Print the shapes of the resulting datasets to verify the split
print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of y_test:", y_test.shape)

```

Shape of X\_train: (1069, 1)  
 Shape of X\_test: (268, 1)  
 Shape of y\_train: (1069,)  
 Shape of y\_test: (268,)

```

[51]: # Get the total number of rows in the original DataFrame
total_rows_df = df.shape[0]
print(f"Total rows in original DataFrame: {total_rows_df}")

# Get the number of rows in the training and testing sets
total_rows_train = X_train.shape[0]
total_rows_test = X_test.shape[0]
print(f"Rows in training set: {total_rows_train}")
print(f"Rows in testing set: {total_rows_test}")

# Sum the rows from training and testing sets
combined_rows = total_rows_train + total_rows_test
print(f"Combined rows (training + testing): {combined_rows}")

# Verify if the combined rows match the original DataFrame rows
if total_rows_df == combined_rows:
    print("All data from the original dataset has been used in the training and",
        ↪testing sets.")
else:
    print("There is a mismatch in the number of rows.")

```

Total rows in original DataFrame: 1337  
 Rows in training set: 1069  
 Rows in testing set: 268  
 Combined rows (training + testing): 1337  
 All data from the original dataset has been used in the training and testing sets.

```

[52]: from sklearn.linear_model import LinearRegression

# Initialize the Linear Regression model
model = LinearRegression()
# Train the model using the training data

```

```

model.fit(X_train, y_train)

print("Linear Regression model trained successfully.")

```

Linear Regression model trained successfully.

```

[53]: # Make predictions on the test set using the trained model
y_pred = model.predict(X_test)

from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

# Calculate R-squared, Mean Absolute Error (MAE), and Mean Squared Error (MSE)
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

# Print the evaluation metrics
print(f"R-squared: {r2:.4f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Error (MSE): {mse:.2f}")

```

R-squared: 0.0951  
 Mean Absolute Error (MAE): 9657.79  
 Mean Squared Error (MSE): 166275348.23

```

[54]: import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

# Set up plot style for better aesthetics
sns.set_style("whitegrid")

# 1. Scatter plot of data with regression line
plt.figure(figsize=(12, 6))
# Plot actual charges against age
plt.scatter(X_test, y_test, alpha=0.6, label='Actual Charges')
# Plot the regression line (predicted charges against age)
plt.plot(X_test, y_pred, color='red', label='Regression Line')
plt.xlabel('Age')
plt.ylabel('Charges')
plt.title('Age vs. Charges with Regression Line')
plt.legend()
plt.show()

# 2. Residuals plot
# Calculate the residuals (actual values - predicted values)
residuals = y_test - y_pred

```

```

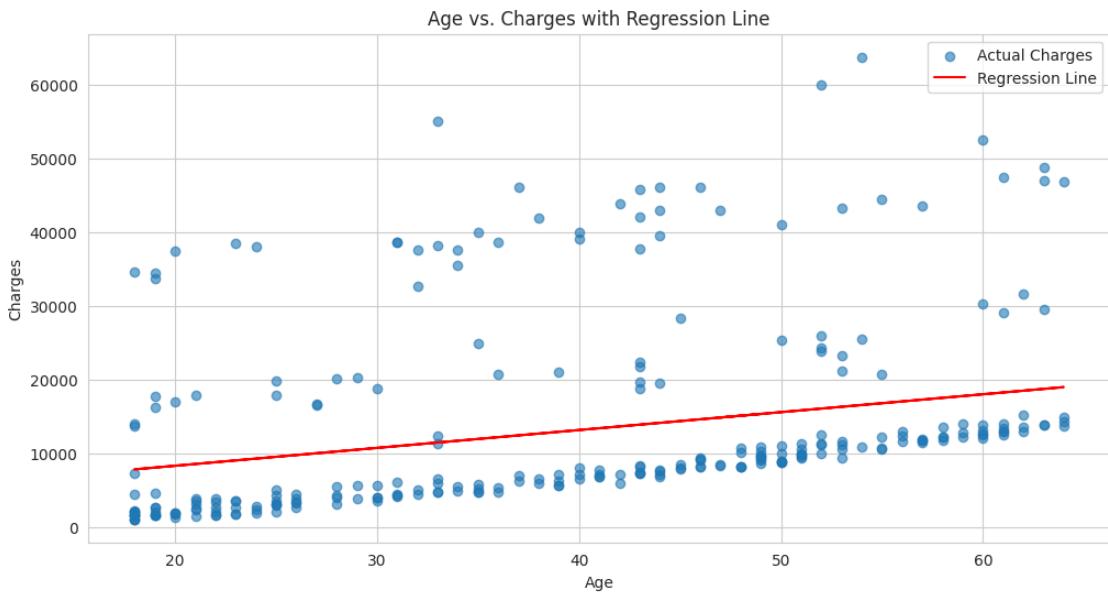
plt.figure(figsize=(12, 6))
# Plot residuals against the independent variable (age)
plt.scatter(X_test, residuals, alpha=0.6)
# Draw a horizontal line at y=0 to easily visualize deviations
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel('Age')
plt.ylabel('Residuals')
plt.title('Residuals Plot (Age vs. Residuals)')
plt.show()

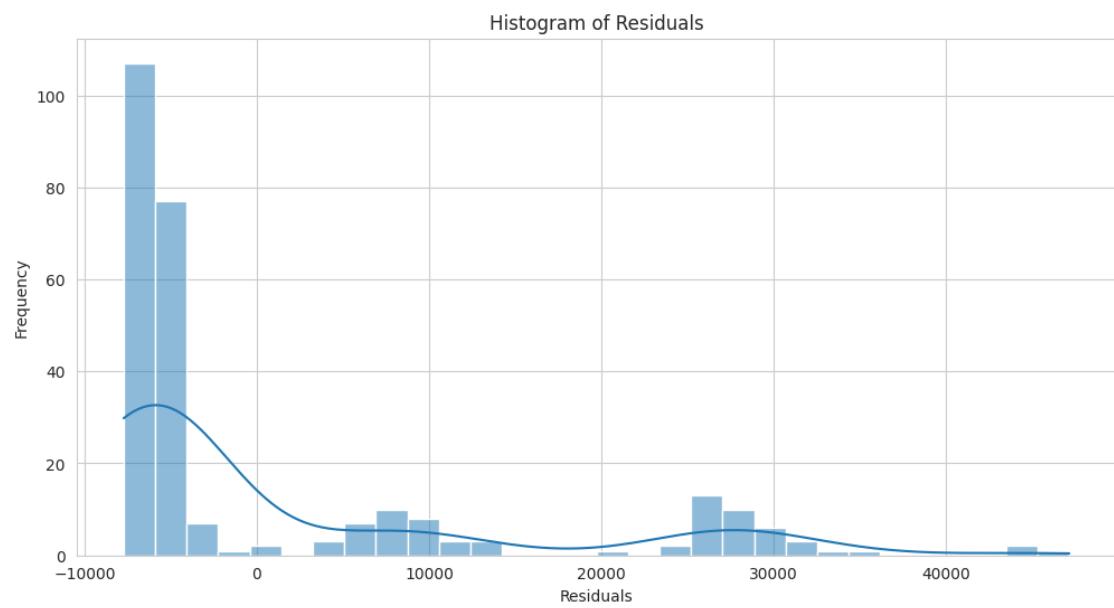
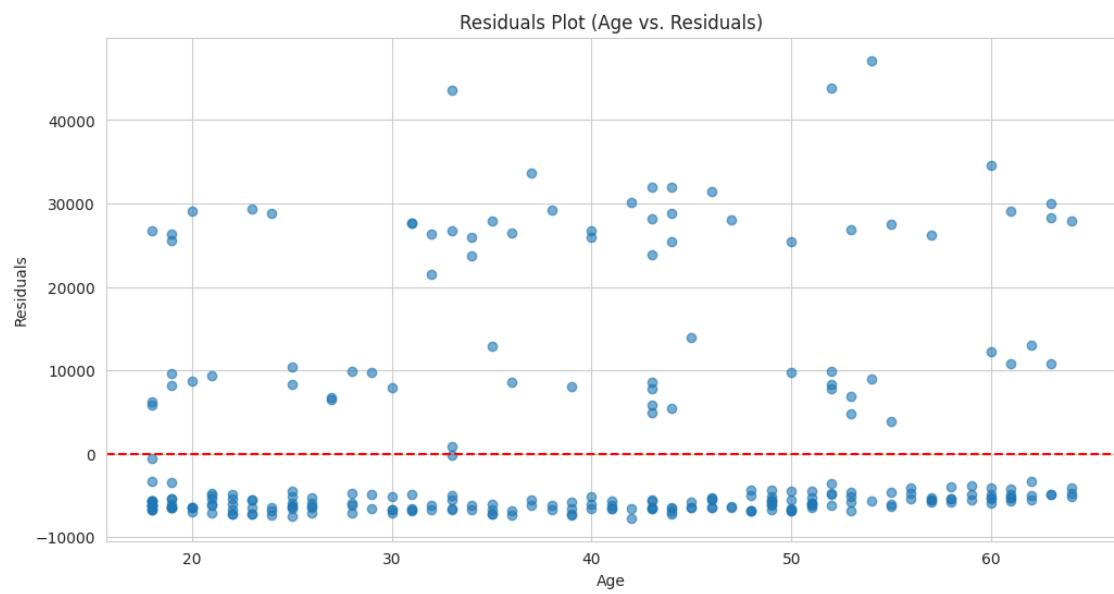
# 3. Histogram of residuals
plt.figure(figsize=(12, 6))
# Plot a histogram of the residuals to check their distribution
sns.histplot(residuals, kde=True, bins=30)
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.title('Histogram of Residuals')
plt.show()

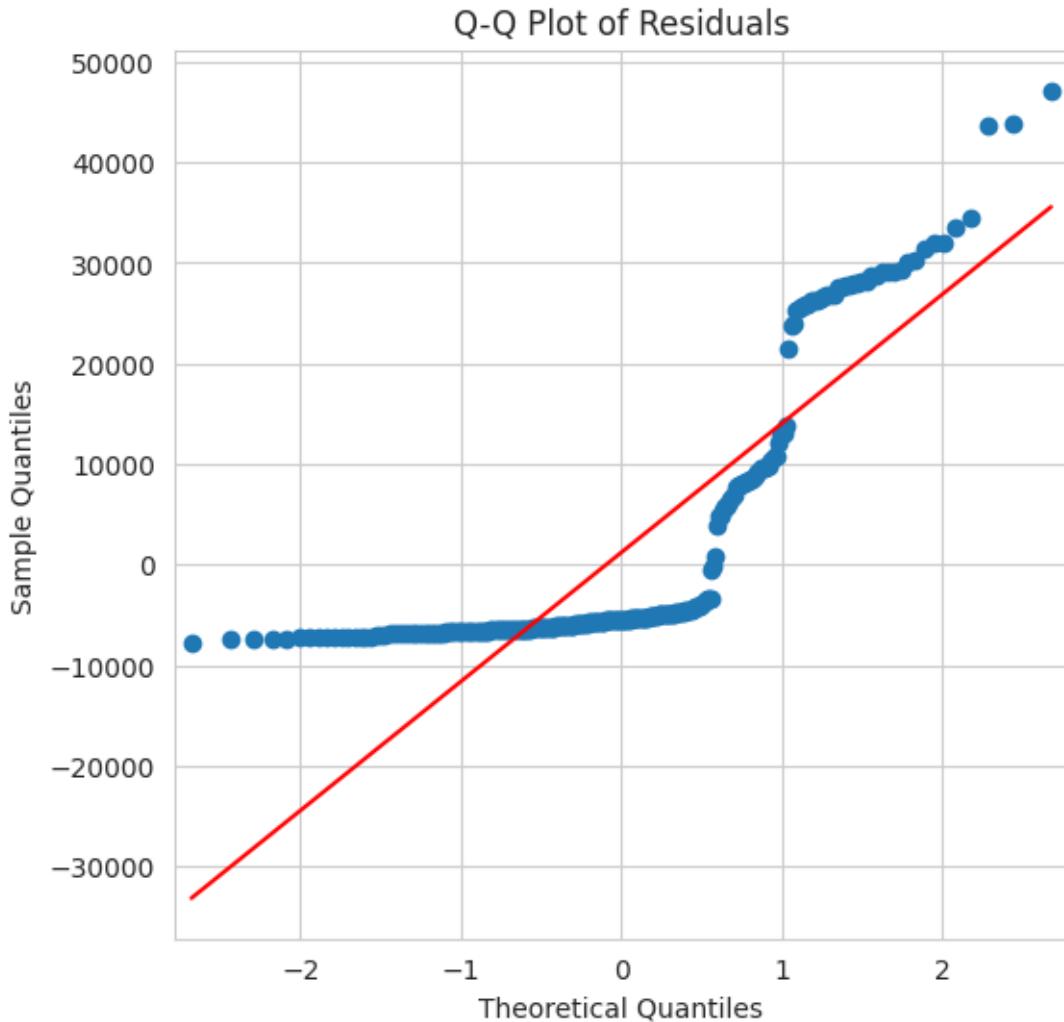
# 4. Q-Q plot of residuals
plt.figure(figsize=(6, 6))
# Create a Q-Q plot to assess if residuals are normally distributed
sm.qqplot(residuals, line='s', ax=plt.gca())
plt.title('Q-Q Plot of Residuals')
plt.show()

print("Visualization of linear regression results completed.")

```







Visualization of linear regression results completed.

## 1.1 Feature Engineering and Selection

```
[55]: import pandas as pd

# Select the specified features from the DataFrame
X_features = df[['age', 'bmi', 'children', 'sex', 'smoker', 'region']].copy()

# Apply one-hot encoding to categorical columns and drop the original columns
X_features = pd.get_dummies(X_features, columns=['sex', 'smoker', 'region'],
                             drop_first=False)

# Display the first few rows of the encoded DataFrame and its info to verify
print("First 5 rows of X_features after encoding:")
```

```

print(X_features.head())
print("\nDataFrame Info for X_features:")
X_features.info()

First 5 rows of X_features after encoding:
   age      bmi  children  sex_female  sex_male  smoker_no  smoker_yes \
0   19  27.900         0       True    False    False      True
1   18  33.770         1      False     True    True    False
2   28  33.000         3      False     True    True    False
3   33  22.705         0      False     True    True    False
4   32  28.880         0      False     True    True    False

   region_northeast  region_northwest  region_southeast  region_southwest
0            False           False           False            True
1            False           False           True            False
2            False           False           True            False
3            False           True           False            False
4            False           True           False            False

DataFrame Info for X_features:
<class 'pandas.core.frame.DataFrame'>
Index: 1337 entries, 0 to 1337
Data columns (total 11 columns):
 #  Column          Non-Null Count  Dtype  
---  --  
 0  age             1337 non-null   int64  
 1  bmi             1337 non-null   float64 
 2  children        1337 non-null   int64  
 3  sex_female      1337 non-null   bool   
 4  sex_male        1337 non-null   bool   
 5  smoker_no       1337 non-null   bool   
 6  smoker_yes      1337 non-null   bool   
 7  region_northeast 1337 non-null   bool  
 8  region_northwest 1337 non-null   bool  
 9  region_southeast 1337 non-null   bool  
 10 region_southwest 1337 non-null   bool  
dtypes: bool(8), float64(1), int64(2)
memory usage: 52.2 KB

```

```
[56]: from sklearn.model_selection import train_test_split

y = df['charges']

# Split the encoded data into training and testing sets (80% train, 20% test)
X_train_multi, X_test_multi, y_train_multi, y_test_multi = train_test_split(X_features, y, test_size=0.2, random_state=42)
```

```
# Print the shapes of the resulting datasets to verify the split
print("Shape of X_train_multi:", X_train_multi.shape)
print("Shape of X_test_multi:", X_test_multi.shape)
print("Shape of y_train_multi:", y_train_multi.shape)
print("Shape of y_test_multi:", y_test_multi.shape)
```

Shape of X\_train\_multi: (1069, 11)  
 Shape of X\_test\_multi: (268, 11)  
 Shape of y\_train\_multi: (1069,)  
 Shape of y\_test\_multi: (268,)

## 1.2 Train Multiple Linear Regression Model

```
[57]: from sklearn.linear_model import LinearRegression

# Initialize the Linear Regression model for multiple features
model_multi = LinearRegression()

# Train the model using the prepared training data
model_multi.fit(X_train_multi, y_train_multi)

print("Multiple Linear Regression model trained successfully.")
```

Multiple Linear Regression model trained successfully.

## 1.3 Evaluate Multiple Linear Regression Model

```
[58]: from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

# Make predictions on the test set using the trained multiple linear regression
# model
y_pred_multi = model_multi.predict(X_test_multi)

# Calculate R-squared, Mean Absolute Error (MAE), and Mean Squared Error (MSE)
r2_multi = r2_score(y_test_multi, y_pred_multi)
mae_multi = mean_absolute_error(y_test_multi, y_pred_multi)
mse_multi = mean_squared_error(y_test_multi, y_pred_multi)

# Print the evaluation metrics
print(f"R-squared (Multiple Linear Regression): {r2_multi:.4f}")
print(f"Mean Absolute Error (MAE) (Multiple Linear Regression): {mae_multi:..2f}")
print(f"Mean Squared Error (MSE) (Multiple Linear Regression): {mse_multi:.2f}")
```

R-squared (Multiple Linear Regression): 0.8069  
 Mean Absolute Error (MAE) (Multiple Linear Regression): 4177.05  
 Mean Squared Error (MSE) (Multiple Linear Regression): 35478020.68

## 1.4 Visualize Multiple Linear Regression Results

```
[59]: import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

# Set up plot style for better aesthetics
sns.set_style("whitegrid")

# 1. Scatter plot of actual vs. predicted charges
plt.figure(figsize=(10, 8))
plt.scatter(y_test_multi, y_pred_multi, alpha=0.6)
plt.plot([min(y_test_multi), max(y_test_multi)], [min(y_test_multi), max(y_test_multi)], 'r--', label='Perfect Prediction Line')
plt.xlabel('Actual Charges')
plt.ylabel('Predicted Charges')
plt.title('Actual vs. Predicted Charges (Multiple Linear Regression)')
plt.legend()
plt.show()

# 2. Residuals plot
# Calculate the residuals (actual values - predicted values)
residuals_multi = y_test_multi - y_pred_multi

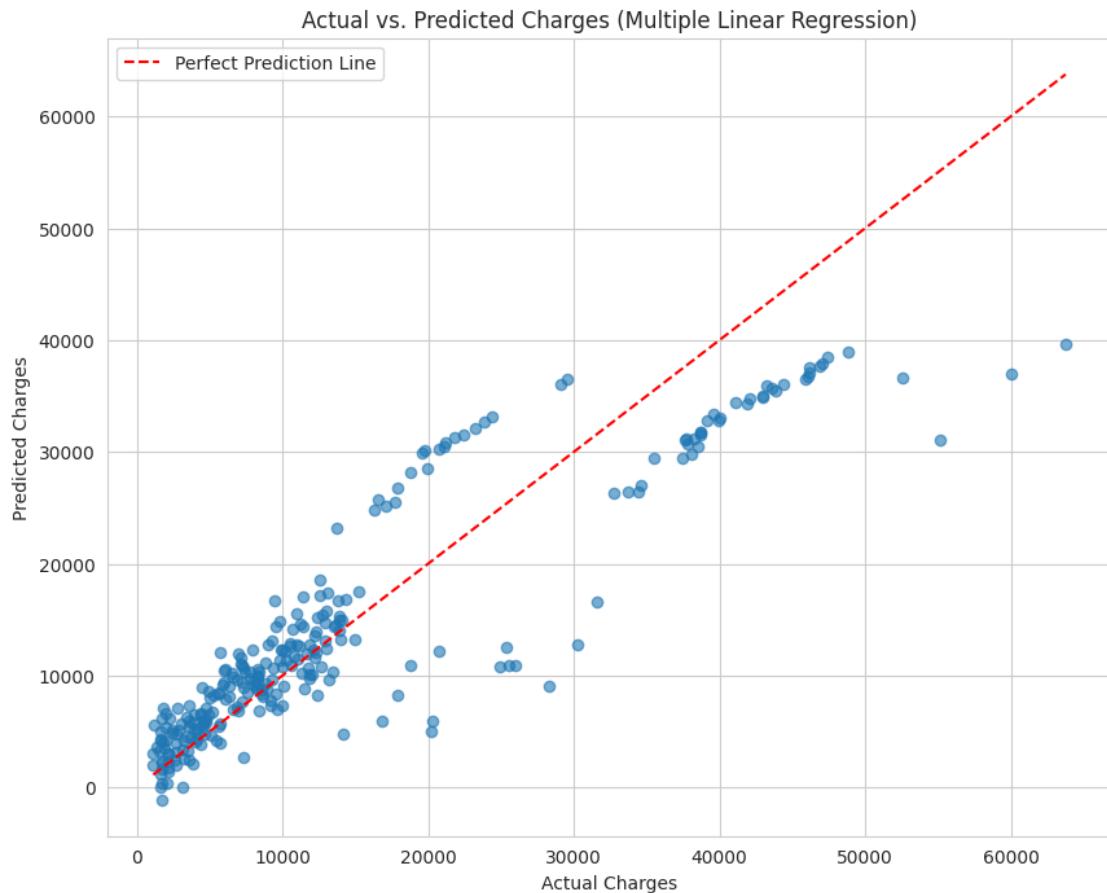
plt.figure(figsize=(10, 6))
# Plot predicted values against residuals
plt.scatter(y_pred_multi, residuals_multi, alpha=0.6)
# Draw a horizontal line at y=0 to easily visualize deviations
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel('Predicted Charges')
plt.ylabel('Residuals')
plt.title('Residuals Plot (Multiple Linear Regression)')
plt.show()

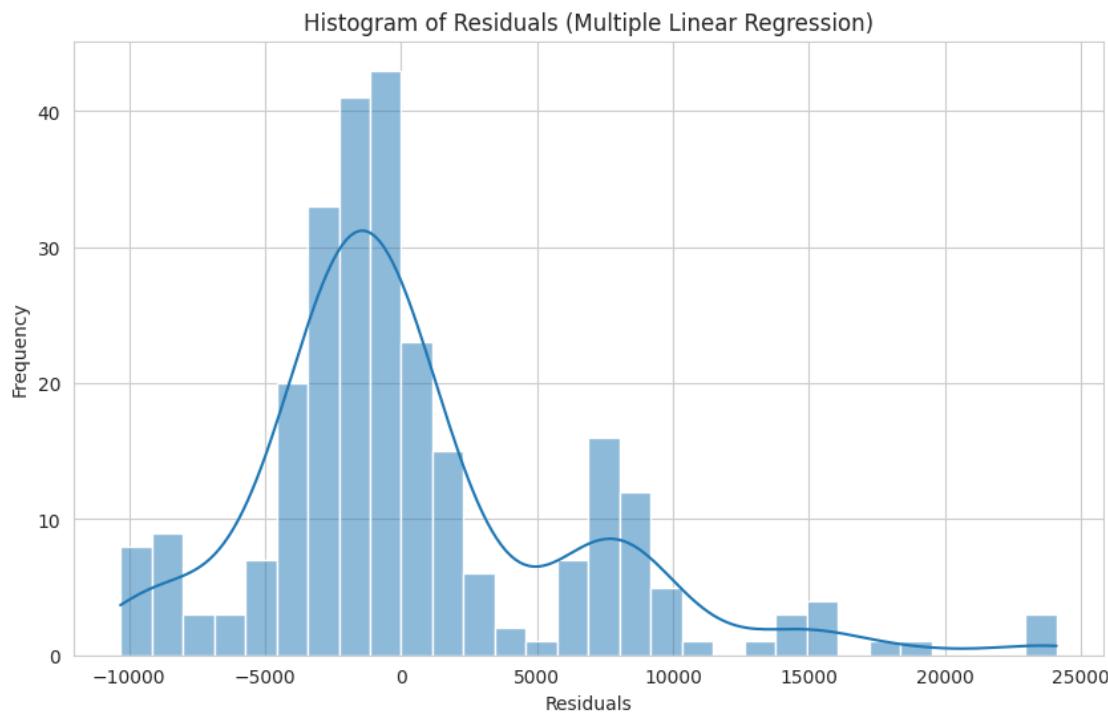
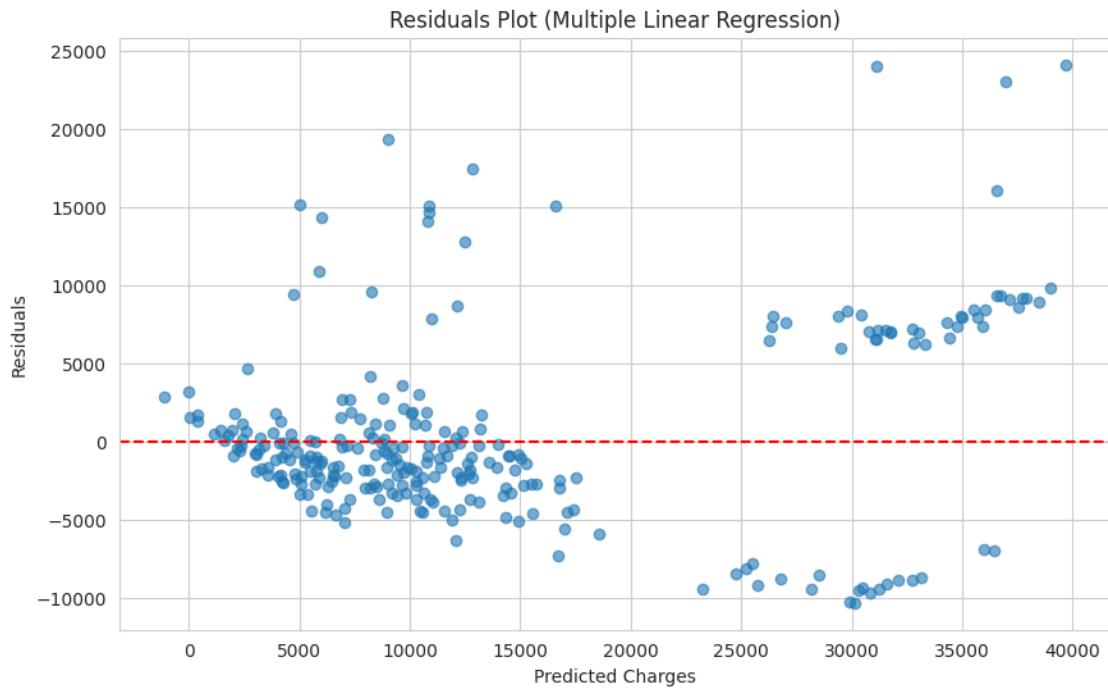
# 3. Histogram of residuals
plt.figure(figsize=(10, 6))
# Plot a histogram of the residuals to check their distribution
sns.histplot(residuals_multi, kde=True, bins=30)
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.title('Histogram of Residuals (Multiple Linear Regression)')
plt.show()

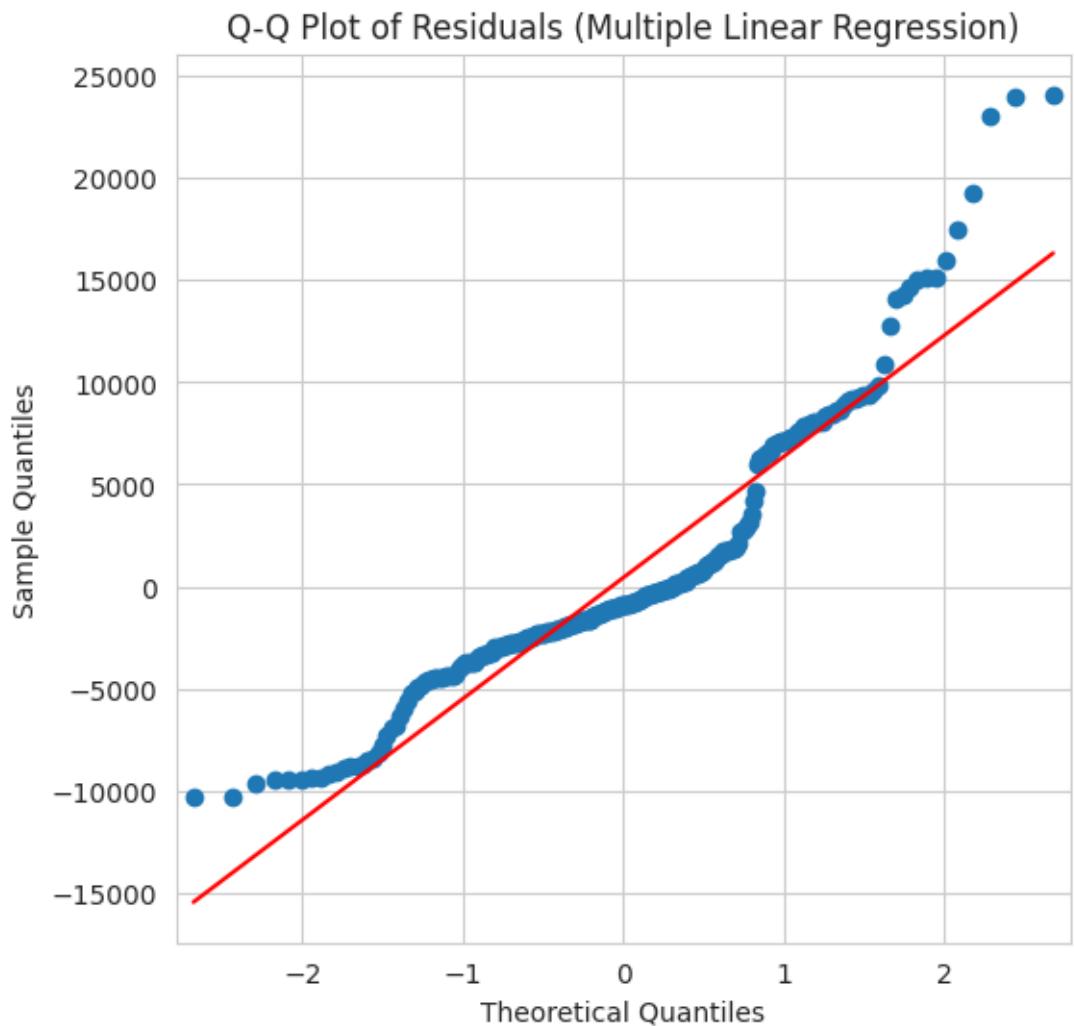
# 4. Q-Q plot of residuals
plt.figure(figsize=(6, 6))
# Create a Q-Q plot to assess if residuals are normally distributed
sm.qqplot(residuals_multi, line='s', ax=plt.gca())
```

```
plt.title('Q-Q Plot of Residuals (Multiple Linear Regression)')
plt.show()

print("Visualization of multiple linear regression results completed.")
```







Visualization of multiple linear regression results completed.

## 1.5 Data Transformation

```
[60]: import numpy as np

# Apply a natural logarithm transformation to the 'charges' column
df['charges_log'] = np.log(df['charges'])

# Display the first few rows of the DataFrame to verify the new 'charges_log' column
print("DataFrame with 'charges_log' column:")
df.head()
```

DataFrame with 'charges\_log' column:

```
[60]:   age      sex      bmi  children smoker      region      charges  charges_log
  0    19  female  27.900        0     yes  southwest  16884.92400    9.734176
  1    18    male  33.770        1    no  southeast   1725.55230    7.453302
  2    28    male  33.000        3    no  southeast   4449.46200    8.400538
  3    33    male  22.705        0    no  northwest  21984.47061   9.998092
  4    32    male  28.880        0    no  northwest  3866.85520    8.260197
```

## 1.6 Advanced Feature Engineering

```
[61]: X_features['smoker_age_interaction'] = X_features['age'] * X_features['smoker_yes']
X_features['smoker_bmi_interaction'] = X_features['bmi'] * X_features['smoker_yes']

# Display the first few rows of the updated DataFrame to verify the new columns
print("X_features DataFrame with interaction terms:")
print(X_features.head())

print("\nDataFrame Info for X_features after adding interaction terms:")
X_features.info()
```

```
X_features DataFrame with interaction terms:
   age      bmi  children  sex_female  sex_male  smoker_no  smoker_yes \
0    19  27.900        0       True     False     False      True
1    18  33.770        1      False      True     True     False
2    28  33.000        3      False      True     True     False
3    33  22.705        0      False      True     True     False
4    32  28.880        0      False      True     True     False

   region_northeast  region_northwest  region_southeast  region_southwest \
0            False          False          False           True
1            False          False          True           False
2            False          False          True           False
3            False          True          False           False
4            False          True          False           False

   smoker_age_interaction  smoker_bmi_interaction
0                      19                  27.9
1                      0                  0.0
2                      0                  0.0
3                      0                  0.0
4                      0                  0.0
```

```
DataFrame Info for X_features after adding interaction terms:
<class 'pandas.core.frame.DataFrame'>
Index: 1337 entries, 0 to 1337
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	age	1337 non-null	int64
1	bmi	1337 non-null	float64
2	children	1337 non-null	int64
3	sex_female	1337 non-null	bool
4	sex_male	1337 non-null	bool
5	smoker_no	1337 non-null	bool
6	smoker_yes	1337 non-null	bool
7	region_northeast	1337 non-null	bool
8	region_northwest	1337 non-null	bool
9	region_southeast	1337 non-null	bool
10	region_southwest	1337 non-null	bool
11	smoker_age_interaction	1337 non-null	int64
12	smoker_bmi_interaction	1337 non-null	float64

dtypes: bool(8), float64(2), int64(3)  
memory usage: 73.1 KB

## 1.7 Prepare Transformed Data

```
[62]: from sklearn.model_selection import train_test_split

# 1. Set the independent variables X_enhanced to the X_features DataFrame
X_enhanced = X_features

# 2. Set the dependent variable y_log to the charges_log column from the
#    original df DataFrame
y_log = df['charges_log']

# 3. Split X_enhanced and y_log into training and testing sets
X_train_enhanced, X_test_enhanced, y_train_log, y_test_log =
    train_test_split(X_enhanced, y_log, test_size=0.2, random_state=42)

# 4. Print the shapes of the new training and testing sets
print("Shape of X_train_enhanced:", X_train_enhanced.shape)
print("Shape of X_test_enhanced:", X_test_enhanced.shape)
print("Shape of y_train_log:", y_train_log.shape)
print("Shape of y_test_log:", y_test_log.shape)
```

Shape of X\_train\_enhanced: (1069, 13)  
Shape of X\_test\_enhanced: (268, 13)  
Shape of y\_train\_log: (1069,)  
Shape of y\_test\_log: (268,)

## 1.8 Train and Evaluate Enhanced Linear Regression Model

```
[63]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
import numpy as np

# Initialize the Linear Regression model for enhanced features and ↴
# log-transformed target
model_enhanced = LinearRegression()

# Train the model using the prepared training data
model_enhanced.fit(X_train_enhanced, y_train_log)

# Make predictions on the test set
y_pred_enhanced = model_enhanced.predict(X_test_enhanced)

# Calculate evaluation metrics
r2_enhanced = r2_score(y_test_log, y_pred_enhanced)
mae_enhanced = mean_absolute_error(y_test_log, y_pred_enhanced)
mse_enhanced = mean_squared_error(y_test_log, y_pred_enhanced)

# Print the evaluation metrics
print(f"R-squared (Enhanced Model with Log-transformed Target): {r2_enhanced:.4f}")
print(f"Mean Absolute Error (MAE) (Enhanced Model with Log-transformed Target): {mae_enhanced:.4f}")
print(f"Mean Squared Error (MSE) (Enhanced Model with Log-transformed Target): {mse_enhanced:.4f}")

# Optionally, transform back to original scale for MAE/MSE if comparison is ↴
# needed with original model
# Note: This is an approximation as np.exp(E[log(y)]) != E[y]
y_pred_original_scale = np.exp(y_pred_enhanced)
y_test_original_scale = np.exp(y_test_log)

mae_original_scale = mean_absolute_error(y_test_original_scale, ↴
y_pred_original_scale)
mse_original_scale = mean_squared_error(y_test_original_scale, ↴
y_pred_original_scale)

print(f"\nMean Absolute Error (MAE) (Original Scale, for comparison): {mae_original_scale:.2f}")
print(f"Mean Squared Error (MSE) (Original Scale, for comparison): {mse_original_scale:.2f}")
```

R-squared (Enhanced Model with Log-transformed Target): 0.8709  
Mean Absolute Error (MAE) (Enhanced Model with Log-transformed Target): 0.2055

Mean Squared Error (MSE) (Enhanced Model with Log-transformed Target): 0.1198

Mean Absolute Error (MAE) (Original Scale, for comparison): 3023.94

Mean Squared Error (MSE) (Original Scale, for comparison): 34328892.68

```
[64]: import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

# Set up plot style for better aesthetics
sns.set_style("whitegrid")

# 1. Scatter plot of actual vs. predicted charges (on log scale)
plt.figure(figsize=(10, 8))
plt.scatter(y_test_log, y_pred_enhanced, alpha=0.6)
plt.plot([min(y_test_log), max(y_test_log)], [min(y_test_log), max(y_test_log)], 'r--', label='Perfect Prediction Line')
plt.xlabel('Actual Log-transformed Charges')
plt.ylabel('Predicted Log-transformed Charges')
plt.title('Actual vs. Predicted Log-transformed Charges (Enhanced Model)')
plt.legend()
plt.show()

# 2. Residuals plot
# Calculate the residuals (actual values - predicted values)
residuals_enhanced = y_test_log - y_pred_enhanced

plt.figure(figsize=(10, 6))
# Plot predicted values against residuals
plt.scatter(y_pred_enhanced, residuals_enhanced, alpha=0.6)
# Draw a horizontal line at y=0 to easily visualize deviations
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel('Predicted Log-transformed Charges')
plt.ylabel('Residuals (Log-transformed)')
plt.title('Residuals Plot (Enhanced Model)')
plt.show()

# 3. Histogram of residuals
plt.figure(figsize=(10, 6))
# Plot a histogram of the residuals to check their distribution
sns.histplot(residuals_enhanced, kde=True, bins=30)
plt.xlabel('Residuals (Log-transformed)')
plt.ylabel('Frequency')
plt.title('Histogram of Residuals (Enhanced Model)')
plt.show()

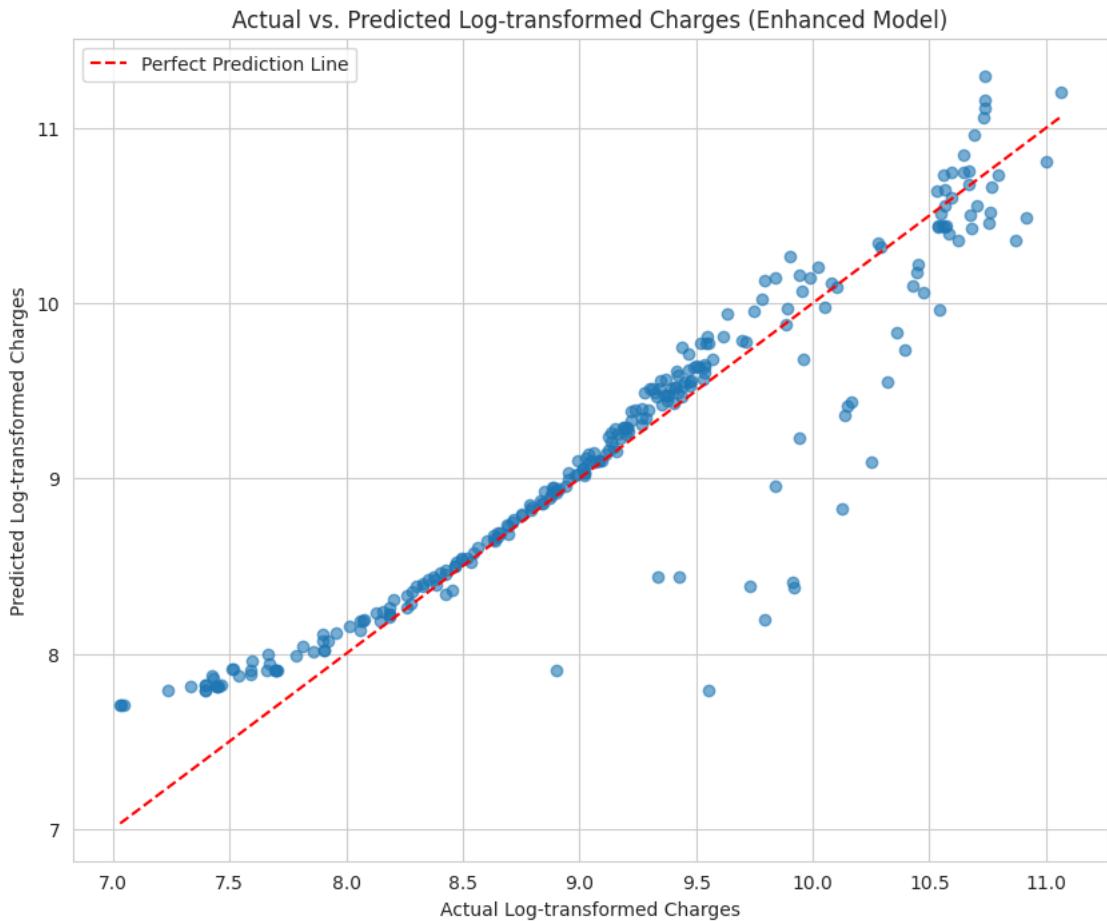
# 4. Q-Q plot of residuals
```

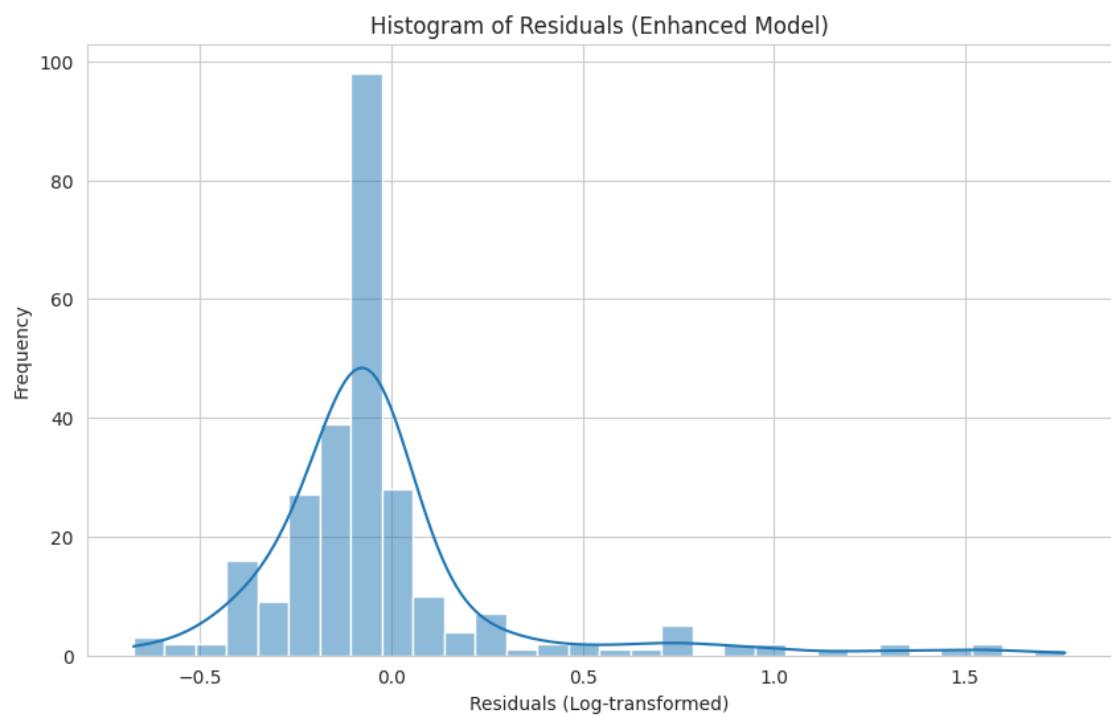
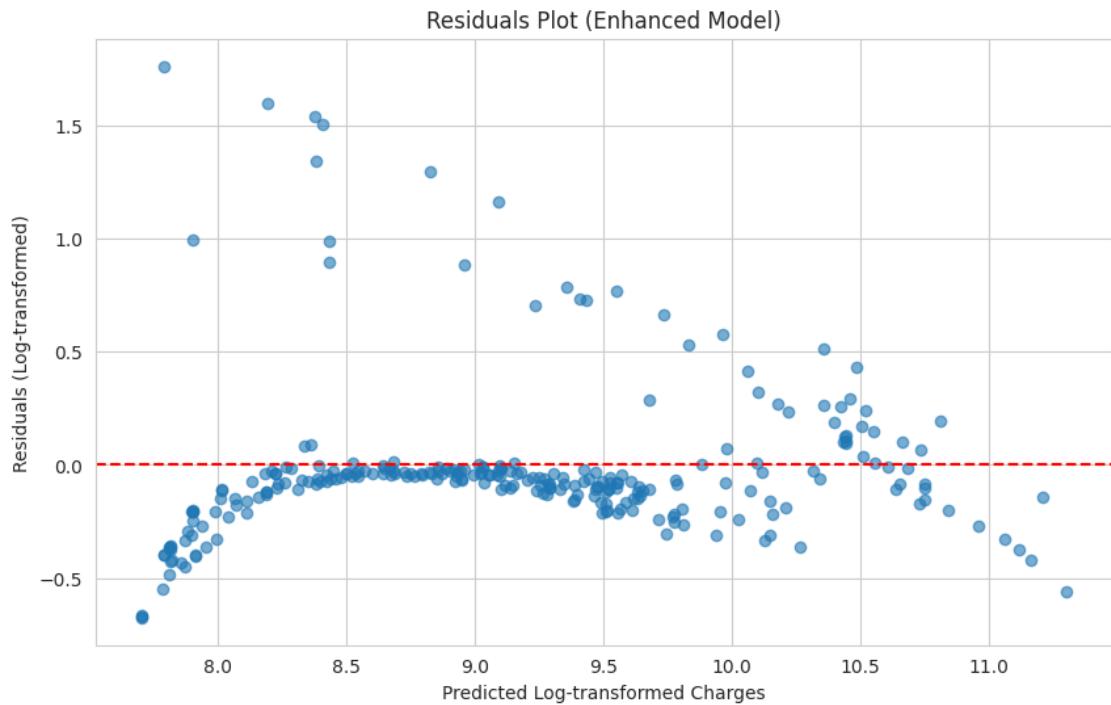
```

plt.figure(figsize=(6, 6))
# Create a Q-Q plot to assess if residuals are normally distributed
sm.qqplot(residuals_enhanced, line='s', ax=plt.gca())
plt.title('Q-Q Plot of Residuals (Enhanced Model)')
plt.show()

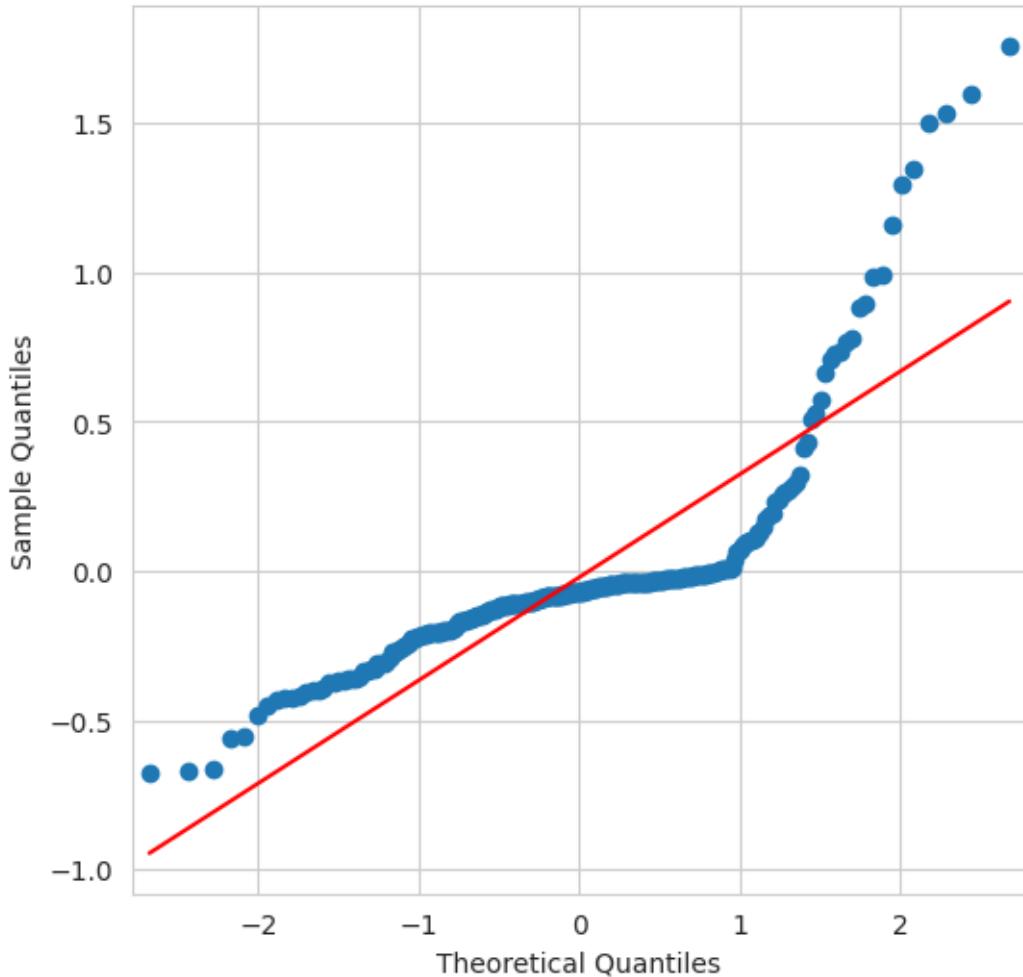
print("Visualization of enhanced linear regression results completed.")

```





Q-Q Plot of Residuals (Enhanced Model)



Visualization of enhanced linear regression results completed.

charges

Data Type: Float (float64) Description: Individual medical costs billed by health insurance (the target variable). Range: \$1121.87 to \$63770.43. Average: Approximately \$13270.42. Standard Deviation: \$12110.01. Insight: The wide range and high standard deviation indicate a highly skewed distribution (many low costs, a few very high costs), making direct prediction challenging without transformations.

Data Quality Missing Values: There are no missing values in any of the columns. Duplicate Entries: One duplicate row was identified and successfully removed, ensuring data uniqueness.

[ ]: